

How to read this

Each tool has two pages. The **Summary page** gives the short version: what the tool is, who has told us what, and where the evidence stands. The **Findings & Evidence** page sets out all fourteen criteria, the score each has been given so far, and what that score rests on. Read the basis column rather than the number: the number is shorthand, the basis is the finding.

Where a score comes from

Every criterion is scored 1 to 4 against wording agreed in advance, so that a 3 means the same thing for every tool. Scores are never added up or averaged, and there is no overall mark. A tool that scores 2 on something your organisation does not care about may still be the right tool for you.

agreed	Reviewers scored it independently, then compared scores and agreed the result on the record. Where they had differed, they settled it by discussion, not by averaging or a vote. None appear in this edition yet.
provisional	Every reviewer who scored it gave the same number, but they have not yet discussed it. None appear in this edition yet.
reviewers differ	Reviewers scored it independently and landed on different numbers. Every score given is shown; the bar on the chart shows the lowest. None appear in this edition yet.
AI-proposed	A draft score with its reasoning, prepared with AI assistance from the evidence we gathered. It is not a human judgement.
contested	The evidence splits, or our own wording does not settle the case. Worth reading the basis.

Where the evidence comes from

Every score carries a tag saying what it rests on. The tags are not interchangeable: what a translator reports from a live project counts for more than what a product page says, and both are printed so you can weigh them yourself.

User-reported	A translator or team using the tool in a live project told us directly, through a survey or an interview we ran.
Facilitator-reported	A facilitator, coordinator or support worker told us about the teams they serve. Close to the work, but second-hand.
Our testing	We used the tool ourselves against a fixed set of questions and recorded what happened.
Documentation	A published page, release note, policy or repository, with the date we read it.
Builder claim	The people who make the tool told us. Printed as theirs, not as established fact.
Pending	Asked for but not yet in hand. The page says what was requested and when.
N/A	The criterion does not apply — for example, explainability for a tool with no AI in it.

The three criteria we treat as decisive

Eleven of the fourteen criteria describe how well a tool does something. Three describe whether it is safe to rely on at all, and a score of 1 on any of them is worth more of your attention than a handful of 2s elsewhere. They are marked in red on the Findings & Evidence page. Nothing here removes a tool from anyone’s list — that decision belongs to your organisation — but these are the three we would want answered before committing a project to a tool.

1. **Problem Alignment** — whether the tool solves a problem translation teams actually have. A tool that is impressive but answers a question nobody asked will not survive contact with a project, however well it scores elsewhere.
11. **Cost & Financial Sustainability** — whether a team can still afford it in three years. Adoption is not a purchase but a dependency: a tool whose funding quietly ends leaves a project stranded mid-book.
12. **Data Governance & Sovereignty** — who holds the translation text, on what terms, and who else can see it. Unpublished Scripture is a community’s own work and sometimes a safety matter, and unlike the other criteria a mistake here cannot be undone later.

The two questions we ask before a tool goes on this list

Can we assess it? That is: is it in scope, and can we find out enough about it? A tool qualifies if it serves Bible translation or related ministry and if we can gather enough information to answer the fourteen questions — from the people using it where they exist, and otherwise from documentation, from the people who built it and from our own use. A tool we cannot find out enough about is named in the wider catalogue and marked as not assessed, rather than being scored on guesswork.

Is it worth your time here? That is: among the tools we could assess, we gave priority to those that address a problem organisations in this room actually have, that more than one team has used, or that stand for a choice many of you are weighing this year. This is a decision about what fits a short session, not a ranking, and a tool being absent says nothing against it.

The fourteen criteria

Every tool is asked the same fourteen questions. ● marks the three we treat as decisive; ● marks the four that need real users before they can be answered. The full wording for each score, 1 to 4, is in the appendix at the back.

1. **Problem Alignment** ● — What problem does it solve?
2. **Setup Complexity** — How easy is it to set up?
3. **Workflow Integration** — Can this tool be easily integrated into the workflow?
4. **User Reach & Adoption** ● — How many people are using it, and who are they?
5. **Ease of Use & Learning Curve** — How intuitive is the tool to use daily?
6. **Project Implementation** ● — How many projects have implemented this tool?
7. **Implementation Horizon** — What can be implemented now, and what is promised later?
8. **Workflow Stage Fit** — Where in its domain's workflow does

the tool fit, and how much of that workflow does it cover?

9. **Post-Implementation Support** ● — What support is available after a team adopts it?
10. **Retention & Sustained Impact** ● — What difference is it making on the ground, over time?
11. **Cost & Financial Sustainability** ● — Is the pricing model viable and sustainable long-term?
12. **Data Governance & Sovereignty** ● — How are data ownership, privacy, and security handled?
13. **Explainability & Human Steering** — Can users understand and control why the AI gives an answer?
14. **Cultural & Safety Alignment** — Are outputs culturally accurate, ethically aligned, and free of regional bias?

Why some pages leave four criteria unscored

Four of the fourteen can only be answered by people using the tool in real work: **4 User Reach & Adoption**, **6 Project Implementation**, **9 Post-Implementation Support** and **10 Retention & Sustained Impact**. Where we could reach such users, a tool is assessed against all fourteen — what we call **Tier 1**. Where we could not, it is assessed against the other ten only, which we call **Tier 2**, and those four are left deliberately unscored rather than guessed at.

A Tier 2 page is not a worse tool or a lesser assessment. It usually means the tool is new, or that its users were beyond our reach in the time we had. Newer tools cannot show a user base until someone adopts them, so treating the absence of one as a mark against a tool would make it impossible for any new tool to be adopted at all. Where a tool is at Tier 2, the page says what we would need in order to answer the remaining four.

Please read this before using the scores. Every score in this document is a draft prepared with AI assistance from the evidence we gathered, and is marked as one. A draft score is not a human judgement. This tool has not yet been scored by our team. The evidence behind every score — the surveys, interviews, correspondence with builders, trial records and our own testing — is the work of our team, and each page says where every finding came from. Please treat the numbers as a starting point for your own judgement, not as settled findings.

P

Paratext Checking Assistant

Paratext · Extension for Paratext 10 · In trial; public release announced for November 2026 ·
Windows, macOS, Linux · Uses AI to check a draft

Executive summary

The Checking Assistant reads a translation draft inside Paratext and puts findings to a reviewer: places where the meaning may have been missed or changed, and places where the language may not sound natural. Each finding appears as a card, and the reviewer decides whether it is a real problem here, a real problem everywhere, or not a problem at all. The assistant never edits the text itself.

This assessment rests on the builder's trial guide, an account from a member of the build team, and the records of the first trial phase: thirteen field reviewers across fifteen projects judged 496 findings between 14 and 28 September, and 54 further points were drawn from their emails. It is the one tool in this document where we can see what reviewers actually did, rather than what they remember doing.

Two findings stand out. Reviewers caught real errors with it, including, one of them reports, in text that had already been through multiple stages of checks. And the trial measured something most trials never reach: how often a reviewer accepted a finding depended more on how many findings they chose to answer than on anything about the checks themselves. Read alongside that, the tool's own severity labels turned out not to predict which findings reviewers accepted.

One limit on this page. The Assistant is announced for public release in November 2026. We assessed the version trialled in September, and we make no claim about what the released product will contain or how it will be installed, priced or supported.

What it is for, and where it fits

Written translation: quality assurance and checking

The second tool in this set at the checking stage, alongside AQuA, and a useful contrast: AQuA scores a whole draft against a reference and shows where the outliers are, while this raises individual findings and asks a reviewer to decide on each

Our connections to this tool

Built by the Paratext team. A member of our evaluation team works on this tool: he gave us the builder's account of it, took part in the trial as a reviewer, and took no part in scoring it.

How much evidence this rests on

Trial records from thirteen reviewers, the builder's guide, and an account from the build team

496

findings judged

13

field reviewers

15

projects

875

findings shown

54

points from emails

0

released users

No released users exist. The records cover 14–28 September 2026; a workshop with further participants falls after this document was finished

Who built it

Built by the Paratext team.

Not established before this document was finished: how the running costs are met once the tool is public, and what the release will contain.

Figures the builder gave us

Supplied by the people who make the tool. We print them as given and have not verified them independently.

FROM THE TRIAL RECORDS

48%

of judged findings were accepted as real

195 of 404 on the model checks. Read on its own this figure is unsafe: more findings were shown and left unanswered than were answered.

FROM THE TRIAL RECORDS

22%

of everything shown was marked real

The floor rather than the rate. The true figure lies between this and 48%, and where depends on why findings were skipped.

FROM THE TRIAL RECORDS


7 of 23

chapters where a suggestion would have made the text worse

The five described in detail share one pattern: the suggested wording was not valid in the target language.

All fourteen questions at a glance

One bar per question, numbered to match the rows overleaf. The taller the bar, the higher the score out of four; the colour shows what kind of evidence that score rests on. Outlined bars are questions we have not been able to answer yet. Hatched bars are draft scores still to be reviewed.



Builder claim

Documentation

User-reported

Pending

AI-proposed

Red

numbers: veto criteria

Where the evidence came from	
User-reported	2
Documentation	9
Builder claim	3

Before scoring

Can we assess it? Yes. In scope, and assessable in unusual depth from the trial records, the written guide and an account from the team. No released users exist, so what we have is trial evidence rather than field evidence.

Is it worth your time here? Yes. It is announced for public release in November 2026, and most people in the room use Paratext. A tool they will be offered within weeks is exactly what an assessment like this should cover, and passing over it would be the odd choice.

What the trial showed

- Reviewers caught real errors with it. One wrote: “‘Disturbing’ because the vast majority of issues raised were legitimate and just mean that the published text the tool was checking is far more imperfect than we believed”; another, writing for the translator who did the review, that “even though the translation text went through multiple stages of checks, the AI identified some translation errors that humans had missed”.
- How often a reviewer accepted a finding depended strongly on how many they chose to answer. Reviewers who answered fewer than two in five accepted 71% of those; reviewers who answered three in five or more accepted 42%. As the trial’s own readout puts it, “reviewers differed less in what they thought of the tool than in how they used the answer buttons.”
- The assistant labels every finding critical, major or minor. Those labels did not predict which findings reviewers accepted — minor findings were accepted at least as often as critical ones. The builder’s own conclusion is not to filter on them.
- Of 23 chapters where reviewers were asked, seven reported a suggested change that would have made the text worse.

The five described in detail all produced wording that is not valid in the target language — a word that does not exist, a gender mismatch, the wrong term for a numeral.

- Not one of 49 project-wide rejections was an acceptance. Either the rules the assistant generalises are poor, or that control is where people go to dismiss a recurring irritation; the trial cannot separate the two.
- Reviewers asked for an answer the interface does not offer: a question that was reasonable but will not change the text. Without it, a useful finding is filed as if it were a false alarm.
- The checks inherited from the older rule-based system were accepted about a quarter of the time. They have been suspended from the next test release while the developers work on improving them.
- Reported gaps: spelling, word order, and the kind of unusual word a human consultant would query. On one chapter the AI provider refused the meaning check outright on content grounds.

Findings & Evidence

Each score sits against a written anchor (1 unsatisfactory, 2 developing, 3 proficient, 4 exemplary). The basis column says what the score rests on and what would change it. Scores marked AI-proposed are drafts prepared from the evidence on file and are still to be reviewed by two people; the appendix at the back gives the wording each score is measured against.

	Criterion	Score	Evidence	Basis
1	Problem Alignment <i>Decisive</i>	3 ■■■■ <i>AI-proposed</i>	Builder claim	The problem is clearly stated: checking is slow and depends on scarce consultant time. The first trial phase gives an unusually direct answer on whether it helps. Thirteen field reviewers across fifteen projects judged 496 findings, and just under half of the findings the models produced were accepted as real. Reviewers describe it catching real errors. One wrote: “‘Disturbing’ because the vast majority of issues raised were legitimate and just mean that the published text the tool was checking is far more imperfect than we believed”; another, writing for the translator who did the review, that “even though the translation text went through multiple stages of checks, the AI identified some translation errors that humans had missed”. Against that, more findings were shown and left unanswered than were answered at all. The tool surfaces real problems; how much of a reviewer’s attention it costs to reach them is the open question.
2	Setup Complexity	2 ■■■■ <i>AI-proposed</i>	Documentation	The trial install is not part of the ordinary installer, and several reviewers lost time to it: a folder that had to exist before the assistant would appear at all, and files left by an earlier beta that stopped the host application starting until they were found and deleted by hand. Once running, reviewers report crashes, failed starts, no recovery after a laptop slept, and an initial analysis that took hours on a whole-Bible project. This describes the trial arrangement rather than the released product, and the team is acting on it: a known-good install with a clean-up step for machines that ran an earlier beta is on their own list for the next phase. Would change it: what installation looks like once it ships
3	Workflow Integration	2 ■■■■ <i>AI-proposed</i>	Documentation	It runs inside the newer version of the host application, beside the text, which is where the work happens. The friction reviewers actually hit is specific and several-sided: a synchronisation step on every switch between the old and new versions, which cannot be turned off; no back-translation view in the newer version; the text of a finding cannot be copied, though consultants want it in their own notes; the guide and the tool give conflicting instructions for exporting a reviewer’s decisions; and earlier results disappear on a re-run. The decisions themselves live only in each machine’s local cache and do not travel with the project.
4	User Reach & Adoption <i>Needs real users</i>	2 ■■■■ <i>AI-proposed</i>	Builder claim	Thirteen field reviewers returned decision ledgers from fifteen projects between 14 and 28 September, joined by two helpdesk reviewers exercising the tool rather than checking a live translation, whose data the team holds apart. Languages include Urdu, Telugu, Limbu and others. This is a controlled trial group rather than a user base, which is what one expects of a tool that has not yet been released.
5	Ease of Use & Learning Curve	2 ■■■■ <i>AI-proposed, contested</i>	Documentation	Divided, and both sides are worth hearing. Several reviewers found it easy to learn, praised the interface and said it exceeded their expectations. Others report that the assistant’s own font is harder to read than the host application’s, that the cards offering language rules do not say what agreeing to one actually does, and that the newer version of the host application does not support Nastaliq — which makes it unusable for Urdu-script languages, and led one reviewer to say they would not recommend it in their region on that ground alone. Would change it: script support settled, and the language-rule cards saying what agreeing to one does

6	Project Implementation <i>Needs real users</i>	2 ■■■■ <i>AI-proposed</i>	Builder claim	Fifteen projects took part in the first trial phase. None adopted the tool in the ordinary sense: they were asked to try it and record what they thought of each finding. There are no projects relying on it, because it is not yet released.
7	Implementation Horizon	3 ■■■■ <i>AI-proposed</i>	Documentation	The trial build is numbered 0.4.1-alpha, public release is announced for November 2026, and the team has said it will spend 2027 examining how the AI checks and the older rule-based checks should relate to each other. The first phase has produced a specific and public list of what to change before the next one. What separates this from the top of the scale is that the release is announced rather than delivered, and we make no claim about what it will contain.
8	Workflow Stage Fit	2 ■■■■ <i>AI-proposed</i>	Documentation	The assistant covers one stage thoroughly: reviewing a chapter for meaning and naturalness and recording a decision on every finding. It hands nothing onward. A reviewer who accepts a finding still makes the edit themselves, and the guide suggests making it in a different application. Decisions stay with the reviewer rather than reaching the rest of the team, and they live only in a local cache that does not travel — so anything not exported before a session ends is lost.
9	Post-Implementation Support <i>Needs real users</i>	3 ■■■■ <i>AI-proposed</i>	Documentation	During the trial, support is named people at the builder, a troubleshooting table covering the faults participants actually hit, and a detailed written guide; helpdesk staff also took part as reviewers, so the people answering questions have used the tool on real text. Participants were invited to send open feedback at any point, and 54 separate points were drawn from their emails. That is more than most trials collect, and it is trial support for a trial group — it tells us nothing about what a released product will provide.
10	Retention & Sustained Impact <i>Needs real users</i>	2 ■■■■ <i>AI-proposed, contested</i>	User-reported	The tool has not been released, so there is nothing to retain anyone yet. What the trial shows is how reviewers behaved once they had it, and the central finding is about them rather than the tool: how often a reviewer accepted a finding depended strongly on how many findings they chose to answer. Reviewers who answered fewer than two in five of the findings they saw accepted 71% of those they answered; reviewers who answered three in five or more accepted 42%. The intervals do not overlap. Someone who presses only on the cards they agree with will show high precision whatever the tool does. Read the other way — accepted out of everything shown — the order reverses. The team's own conclusion is to ask for an answer on every finding in the next phase. Would change it: a phase where every finding gets an answer, so precision and acceptance-per-finding-shown become the same number
11	Cost & Financial Sustainability <i>Decisive</i>	3 ■■■■ <i>AI-proposed, contested</i>	Documentation	If the Checking Assistant disappeared in two years, a reviewer's own decisions and any notes they had written would remain in the Paratext project, because the tool never holds the work: it reads the text and puts findings to a person. What would be lost is the flow of findings itself, and a team that had built its checking round them would go back to checking as before. Nothing about the tool creates a cost for a team: it is free to trial participants, it is an extension to software they already run, and there is no licence, subscription or per-seat charge. The dependency sits with the builder rather than the user — every check is a paid call to an outside AI service, so the tool costs its builder money each time anyone uses it, and the trial build already caps how many checks a session allows. Who meets that cost once the tool is public, and whether the cap survives, is not settled, because the tool is not yet released and its commercial arrangements do not yet exist. That is a fact about a pre-release product rather than a gap in our research, and it is the first thing an adopting organisation should ask when the release is announced. Would change it: a stated arrangement

				for the running costs after release. A scorer who reads an unsettled post-release position as reliance on an unsustainable subsidy would score lower; one who reads free-to-the-team with costs carried elsewhere as the top of the scale would score higher. The case for 3 is that nothing about the tool bills a team, and nothing suggests strain — but the arrangement that will apply does not exist yet.
12	Data Governance & Sovereignty Decisive	3 ■■■■ <i>AI-proposed</i>	Documentation	The trial guide gives a full account of how data is handled. It names the AI service used, states that the provider does not use the data to train its models, explains why that provider was chosen and links to the specific legal terms. The assistant reads the translation and never edits it, and the host application works on a copy from its server which keeps every revision. The record a participant sends back contains only the few words each finding pointed at, plus anything the reviewer typed. One incident from the trial is worth knowing: on one chapter the AI provider refused the meaning check outright on content grounds, and re-running did not get past it. What is not stated is how long anything is kept or where it is processed.
13	Explainability & Human Steering	3 ■■■■ <i>AI-proposed, contested</i>	Documentation	Strong in its framing, and with one measured weakness the trial found. The guide tells reviewers plainly that the assistant may miss real problems, may raise ones that are not real, can be confidently wrong in the same tone it uses when right, and may produce different results if run twice. Findings are grounded in material the reviewer can see and influence, and rejection can be scoped to one verse or the whole project. The weakness is severity: the assistant marks every finding critical, major or minor, and on the meaning leg those labels do not predict which findings reviewers accepted — minor findings were accepted at least as often as critical ones. A reviewer who trusted the labels to triage their attention would be misled, and the team's own recommendation is not to filter on them. One further gap reviewers named: the interface offers only real or not-an-issue, with no way to say that a question was reasonable but will not change the text, so a useful finding has to be filed as if it were a false alarm. Would change it: severity labels that predict correctness, and a third answer for a reasonable question that changes nothing
14	Cultural & Safety Alignment	2 ■■■■ <i>AI-proposed, contested</i>	User-reported	The assistant checks meaning against a literal open-source reference translation, so that reference carries exegetical choices a project may not share. The trial surfaced something sharper than that, in the tool's suggested edits. Reviewers were asked whether the assistant ever proposed a change that would have made the text worse; of 23 chapters where the question was answered, seven said yes, and the five described in detail share one pattern — the suggestion produced text that is not valid in the target language. A word that does not exist, formed by analogy with the national language. A misspelt insertion carrying the wrong noun. A mismatch of grammatical gender. The wrong term for a numeral. A pronoun altered inside reported thought where the original was right. One reviewer pointed at an obvious guard: check that a word occurs in the project's own text before offering it. Separately, the language-rule checks rest on a profile built from each project's own draft, so they work well where there is text to learn from and badly where there is not — on a near-empty scratch project the same rule that scored 18 out of 18 on real projects flagged ordinary Telugu letters as rare. Would change it: a check that a proposed word occurs in the project's own text before it is offered

Still open

- Whether the next trial phase, which asks for an answer on every finding, closes the gap between the two ways of counting. Until it does, the honest answer on accuracy is a range, not a number.
- Why reviewers press the project-wide button only to reject. Two explanations fit the data and only a live session can separate them.
- Who meets the running costs once the tool is public, and whether the per-session cap survives. Every check is a paid call to an outside service. This is unsettled because the tool is not yet released, and it is the first question to ask when it is.
- What the November release contains, how it is installed and what support comes with it. We deliberately make no claim about any of it.
- How teams work with it together. Decisions stay on one machine and do not travel, so the trial could only show individuals working alone.
- Script support. The newer host application does not handle Nastaliq, which rules out Urdu-script languages until it does.

Builder's note. The builder has not yet reviewed this page.

Where to find it

Not publicly available at the time of writing.

Announced for release in November 2026 through Paratext 10. Trial participants install it as an extension; the builder distributes the guide and the files directly.

The user it assumes

A reviewer or consultant already working in Paratext, with a reliable internet connection, since every check is a call to an outside service. During the trial it also assumes one designated computer, a willingness to install software by hand into a hidden folder, an hour and a half for the initial analysis, and enough confidence to overrule a tool that states its case firmly whether or not it is right.

How to read this

A score is chosen by reading the four descriptions in the appendix and picking the one that fits the evidence.

AI-proposed means a draft score, prepared from the evidence on file and still to be reviewed by two people.

Contested means the evidence points both ways, or the rubric's own wording does not settle the case — the reasoning beside it is worth reading. Figures the builder gave us are labelled as theirs. Scores are never added together or averaged.

How this page was put together. Findings were gathered by the evaluation team from the builder's trial records, participant feedback and written guide, and an account given by a member of the build team. Reviewers and projects are pseudonymised in the source and are not identified here. The scores are drafts prepared with AI assistance from that evidence and are still to be reviewed by two people.

Appendix: the full rubric

The wording every score is measured against. A score is chosen by reading these four descriptions and picking the one that fits what the evidence shows — not by impression, and not on a curve. ● decisive criterion · ● needs real users (Tier 1 only).

Criterion	1 — Unsatisfactory	2 — Developing	3 — Proficient	4 — Exemplary
1 Problem Alignment <i>What problem does it solve?</i> ● decisive	Solves an ill-defined or non-existent problem; no alignment with target audience needs.	Addresses a minor issue, but core pain points remain unresolved.	Solves a clear, well-defined problem relevant to field translation teams.	Solves a high-priority, critical problem with significant, measurable impact.
2 Setup Complexity <i>How easy is it to set up?</i>	Extremely complex; requires specialized technical skills, custom code, or extensive workarounds.	Difficult setup; requires moderate technical support or multi-step configuration.	Straightforward setup; standard user can complete configuration with documentation.	Turnkey, plug-and-play setup; minimal to no technical intervention required.
3 Workflow Integration <i>Can this tool be easily integrated into the workflow?</i>	Incompatible; disrupts existing workflows and creates significant friction or duplicate work.	Integrates with difficulty; requires manual data re-entry or awkward workarounds.	Integrates smoothly into standard field tools and processes with minimal friction.	Seamless native integration; fully automated data flow between tools and platforms.
4 User Reach & Adoption <i>How many people are using it, and who are they?</i> ● needs real users	No reachable active users; purely theoretical or tested only in demo settings.	Very few active users (1–5); restricted to a single closed group.	Moderate active user base across target user roles.	Large, established active user base across multiple field teams and roles.
5 Ease of Use & Learning Curve <i>How intuitive is the tool to use daily?</i>	Unintuitive interface; steep learning curve requiring extensive ongoing training or technical intervention.	Somewhat complex interface; requires moderate user training and frequent reliance on support manuals.	User-friendly and intuitive for everyday translation work; minimal training required for field teams.	Highly intuitive and accessible; seamless navigation with virtually no learning curve for non-technical users.
6 Project Implementation <i>How many projects have implemented this tool?</i> ● needs real users	0 real-world projects implemented.	1 active pilot project.	2–4 active field projects currently using the tool.	5+ real-world translation projects actively relying on the tool.
7 Implementation Horizon <i>What can be implemented now, and what is promised later?</i>	Unclear immediate utility and no long-term roadmap.	Limited current functionality; long-term roadmap exists but lacks timeline or feasibility.	Core features ready for immediate deploy; clear long-term feature roadmap in place.	Fully functional immediate capabilities paired with a realistic, high-value long-term plan.

Criterion	1 — Unsatisfactory	2 — Developing	3 — Proficient	4 — Exemplary
8 Workflow Stage Fit <i>Where in its domain's workflow does the tool fit, and how much of that workflow does it cover?</i>	Cannot be placed at any recognised stage of its declared domain's workflow, or the fit is so narrow or ill-defined that a team could not say when in their process to use it.	Usable at a single stage, but with significant constraints or gaps; the stage's core task is only partly addressed.	Fully functional at one or more core stages of its domain's workflow, with a clear hand-off into and out of the tool.	Covers multiple stages of its domain's workflow with continuity between them — outputs of one stage feed the next without re-entry.
9 Post-Implementation Support <i>What support is available after a team adopts it?</i> ● needs real users	No post-launch support, contact points, or documentation available.	Minimal support; limited to unmaintained community forums or slow response times.	Reliable support available via direct contact channels with clear user documentation.	Comprehensive support including direct helpdesk, ongoing training, and regular updates.
10 Retention & Sustained Impact <i>What difference is it making on the ground, over time?</i> ● needs real users	Tool is abandoned shortly after launch; zero sustained adoption or measurable impact on ground operations.	High drop-off rate after initial trial; minimal lasting change in team productivity or workflows over 6 months.	Consistent usage sustained after 6+ months; clear, positive shift in field team efficiency or output quality.	Entrenched as an essential tool with near-total user retention; measurable, long-term transformational impact on ground work.
11 Cost & Financial Sustainability <i>Is the pricing model viable and sustainable long-term?</i> ● decisive	Cost-prohibitive or opaque pricing; reliant on unsustainable short-term subsidies with high financial risk.	High operational or licensing costs; partially sustainable but creates budget strain over time.	Predictable, fair pricing structure supported by a clear, long-term financial plan.	Free or open-access, or highly cost-effective; backed by a fully sustainable funding and maintenance model.
12 Data Governance & Sovereignty <i>How are data ownership, privacy, and security handled?</i> ● decisive	No data privacy controls; vendor retains unrestricted ownership or rights over sensitive translation content.	Basic security measures, but vague data retention, export, and regional privacy compliance policies.	Clear data ownership rights retained by the translation team, with standard privacy and export protocols.	Complete data sovereignty guaranteed; robust encryption, strict privacy safeguards, and full local data control.
13 Explainability & Human Steering <i>Can users understand and control why the AI gives an answer?</i>	Opaque “black box” system; no rationale provided for outputs, and users cannot adjust, override, or guide model decision-making.	Minimal rationale or source tracking; limited ability to tweak prompt settings or influence underlying output logic.	Clear output explanations or source references provided; users can easily adjust steering parameters and override AI outputs.	Fully transparent logic, traceable citations, and granular human controls to tune, override, and direct AI decision-making at every stage.
14 Cultural & Safety Alignment <i>Are outputs culturally accurate, ethically aligned, and free of regional bias?</i>	Outputs display significant regional bias, offensive or harmful content, or complete disregard for local cultural context.	Inconsistent safety controls; occasional cultural inaccuracies, subtle regional biases, or weak ethical boundary safeguards.	Generally reliable safety and ethical boundaries; culturally appropriate outputs with minimal regional bias.	Rigorously calibrated for localized ethics, nuance, and safety; zero regional bias with complete context awareness and customization.