

How to read this

Each tool has two pages. The **Summary page** gives the short version: what the tool is, who has told us what, and where the evidence stands. The **Findings & Evidence** page sets out all fourteen criteria, the score each has been given so far, and what that score rests on. Read the basis column rather than the number: the number is shorthand, the basis is the finding.

Where a score comes from

Every criterion is scored 1 to 4 against wording agreed in advance, so that a 3 means the same thing for every tool. Scores are never added up or averaged, and there is no overall mark. A tool that scores 2 on something your organisation does not care about may still be the right tool for you.

agreed	Reviewers scored it independently, then compared scores and agreed the result on the record. Where they had differed, they settled it by discussion, not by averaging or a vote. None appear in this edition yet.
provisional	Every reviewer who scored it gave the same number, but they have not yet discussed it. None appear in this edition yet.
reviewers differ	Reviewers scored it independently and landed on different numbers. Every score given is shown; the bar on the chart shows the lowest. None appear in this edition yet.
AI-proposed	A draft score with its reasoning, prepared with AI assistance from the evidence we gathered. It is not a human judgement.
contested	The evidence splits, or our own wording does not settle the case. Worth reading the basis.

Where the evidence comes from

Every score carries a tag saying what it rests on. The tags are not interchangeable: what a translator reports from a live project counts for more than what a product page says, and both are printed so you can weigh them yourself.

User-reported	A translator or team using the tool in a live project told us directly, through a survey or an interview we ran.
Facilitator-reported	A facilitator, coordinator or support worker told us about the teams they serve. Close to the work, but second-hand.
Our testing	We used the tool ourselves against a fixed set of questions and recorded what happened.
Documentation	A published page, release note, policy or repository, with the date we read it.
Builder claim	The people who make the tool told us. Printed as theirs, not as established fact.
Pending	Asked for but not yet in hand. The page says what was requested and when.
N/A	The criterion does not apply — for example, explainability for a tool with no AI in it.

The three criteria we treat as decisive

Eleven of the fourteen criteria describe how well a tool does something. Three describe whether it is safe to rely on at all, and a score of 1 on any of them is worth more of your attention than a handful of 2s elsewhere. They are marked in red on the Findings & Evidence page. Nothing here removes a tool from anyone’s list — that decision belongs to your organisation — but these are the three we would want answered before committing a project to a tool.

1. **Problem Alignment** — whether the tool solves a problem translation teams actually have. A tool that is impressive but answers a question nobody asked will not survive contact with a project, however well it scores elsewhere.
11. **Cost & Financial Sustainability** — whether a team can still afford it in three years. Adoption is not a purchase but a dependency: a tool whose funding quietly ends leaves a project stranded mid-book.
12. **Data Governance & Sovereignty** — who holds the translation text, on what terms, and who else can see it. Unpublished Scripture is a community’s own work and sometimes a safety matter, and unlike the other criteria a mistake here cannot be undone later.

The two questions we ask before a tool goes on this list

Can we assess it? That is: is it in scope, and can we find out enough about it? A tool qualifies if it serves Bible translation or related ministry and if we can gather enough information to answer the fourteen questions — from the people using it where they exist, and otherwise from documentation, from the people who built it and from our own use. A tool we cannot find out enough about is named in the wider catalogue and marked as not assessed, rather than being scored on guesswork.

Is it worth your time here? That is: among the tools we could assess, we gave priority to those that address a problem organisations in this room actually have, that more than one team has used, or that stand for a choice many of you are weighing this year. This is a decision about what fits a short session, not a ranking, and a tool being absent says nothing against it.

The fourteen criteria

Every tool is asked the same fourteen questions. ● marks the three we treat as decisive; ● marks the four that need real users before they can be answered. The full wording for each score, 1 to 4, is in the appendix at the back.

1. **Problem Alignment** ● — What problem does it solve?
2. **Setup Complexity** — How easy is it to set up?
3. **Workflow Integration** — Can this tool be easily integrated into the workflow?
4. **User Reach & Adoption** ● — How many people are using it, and who are they?
5. **Ease of Use & Learning Curve** — How intuitive is the tool to use daily?
6. **Project Implementation** ● — How many projects have implemented this tool?
7. **Implementation Horizon** — What can be implemented now, and what is promised later?
8. **Workflow Stage Fit** — Where in its domain's workflow does

the tool fit, and how much of that workflow does it cover?

9. **Post-Implementation Support** ● — What support is available after a team adopts it?
10. **Retention & Sustained Impact** ● — What difference is it making on the ground, over time?
11. **Cost & Financial Sustainability** ● — Is the pricing model viable and sustainable long-term?
12. **Data Governance & Sovereignty** ● — How are data ownership, privacy, and security handled?
13. **Explainability & Human Steering** — Can users understand and control why the AI gives an answer?
14. **Cultural & Safety Alignment** — Are outputs culturally accurate, ethically aligned, and free of regional bias?

Why some pages leave four criteria unscored

Four of the fourteen can only be answered by people using the tool in real work: **4 User Reach & Adoption**, **6 Project Implementation**, **9 Post-Implementation Support** and **10 Retention & Sustained Impact**. Where we could reach such users, a tool is assessed against all fourteen — what we call **Tier 1**. Where we could not, it is assessed against the other ten only, which we call **Tier 2**, and those four are left deliberately unscored rather than guessed at.

A Tier 2 page is not a worse tool or a lesser assessment. It usually means the tool is new, or that its users were beyond our reach in the time we had. Newer tools cannot show a user base until someone adopts them, so treating the absence of one as a mark against a tool would make it impossible for any new tool to be adopted at all. Where a tool is at Tier 2, the page says what we would need in order to answer the remaining four.

Please read this before using the scores. Every score in this document is a draft prepared with AI assistance from the evidence we gathered, and is marked as one. A draft score is not a human judgement. This tool has not yet been scored by our team. The evidence behind every score — the surveys, interviews, correspondence with builders, trial records and our own testing — is the work of our team, and each page says where every finding came from. Please treat the numbers as a starting point for your own judgement, not as settled findings.



Scribe

Bridge Connectivity Solutions · [scribe.bible](#) · Free, open source · Windows, Linux, macOS and web · No AI

Executive summary

Scribe is an offline editor for oral Bible translation. A team records a passage verse by verse on a laptop, with reference material beside them and the interface in a language they read, then checks and re-records within the same tool. Everything stays on the machine unless the team chooses to sync it.

This assessment rests on a structured interview with the builder on 22 September 2026, the tool’s public documentation and its code repository. We have not yet spoken to anyone who uses it, so three of the fourteen questions are unanswered and several others rest on the builder’s own account.

More than forty teams have drafted with it and forty oral Bibles have been published from those drafts, all within the builder’s own translation programme. The tool reached version 1.2.0 in July 2026 and the builder describes it as finished and maintained rather than still growing. A call is being arranged with the facilitator who supports these teams, and team interviews depend on what comes of it; language and distance make direct contact difficult.

What it is for, and where it fits

Oral translation: drafting and peer checking

It has also been tried for sign language in the Solomon Islands, which we have named but not assessed because those teams cannot be reached. Its text mode has no users we know of

How much evidence this rests on

Builder interview and documentation; no users reached yet

40+

teams have drafted with it

40

oral Bibles published

0

users interviewed yet

3

questions unanswered

Reachable users exist — more than forty teams in the builder’s clusters — and a call with the facilitator who supports them is being arranged

Our connections to this tool

Scribe is built by Bridge Connectivity Solutions, which employs a member of our evaluation team and is a sponsor of this summit. He gave us the builder’s account of the tool and took no part in scoring it; the two people who will score it have no connection to the company.

Who built it

Built by Bridge Connectivity Solutions. What this page says about how the tool works comes mainly from the builder.

Not known: how many teams are working in it now, and how far ahead its upkeep is secured.

Figures the builder gave us

Supplied by the people who make the tool. We print them as given and have not verified them independently.

BUILDER-REPORTED

40

oral Bibles published from Scribe drafts

Across several parallel clusters in the builder's own programme. The published audio is a studio recording made from the checked text, not the recording made in Scribe.

BUILDER-REPORTED

40+

teams have drafted with it

A historical count. How many are working in it today is not known.

NOT KNOWN


—

teams active now

How many of the teams that have drafted with it are still working in it today.

All fourteen questions at a glance

One bar per question, numbered to match the rows overleaf. The taller the bar, the higher the score out of four; the colour shows what kind of evidence that score rests on. Outlined bars are questions we have not been able to answer yet; dotted bars are questions that do not apply to this tool. Hatched bars are draft scores still to be reviewed.



Where the evidence came from	
Documentation	1
Builder claim	9
Pending	3
Not applicable	1

Builder claim

Documentation

Not applicable

Pending

AI-proposed

Red

numbers: veto criteria

Before scoring

Can we assess it? Yes. In scope; assessable from documentation, the public repository and a builder interview. Reachable users exist in BCS's clusters but have not been interviewed, so Tier 1 pending.

Is it worth your time here? Yes. "Which editor for oral teams" has few established answers; evidence exists from many teams, though all within one programme; OBT tooling is a live decision for organisations in the room.

What the builder told us

- It is built to work without internet, with the interface in languages the translators read. The builder lists English, Hindi, Telugu, Kannada, French and Russian; the documentation lists a slightly different set, and we have recorded both.
- Teams record verse by verse and can combine or split recordings. Peer checking and re-recording happen inside Scribe, with a history of revisions kept.
- Finished recordings leave as Scripture Burrito files and go to the builder's companion app, where they are transcribed, checked as text and turned back into speech for review. The round trip keeps verse alignment, and the pipeline passes through text at the checking stage.
- Version 1.2.0 was released on 2 July 2026. The builder describes the tool as complete and maintained, with no new features planned.
- Data stays on the team's device unless they sync it, either to a shared repository, which needs an account and acceptance of an open licence, or to a server they run themselves.

Findings & Evidence

Each score sits against a written anchor (1 unsatisfactory, 2 developing, 3 proficient, 4 exemplary). The basis column says what the score rests on and what would change it. Scores marked AI-proposed are drafts prepared from the evidence on file and are still to be reviewed by two people; the appendix at the back gives the wording each score is measured against.

	Criterion	Score	Evidence	Basis
1	Problem Alignment <i>Decisive</i>	3 ■■■■ <i>AI-proposed</i>	Builder claim	The builder describes a clear problem: oral translation teams working without reliable internet need to record verse by verse, in an interface they can read, with reference material beside them. That is a real and well-defined need. To go higher we would need teams themselves to say it solved the problem they adopted it for.
2	Setup Complexity	— ■■■■	Pending	We do not yet know who installs it, whether teams use the desktop or the web version, or how long it takes to get to a first recording. A call with the facilitator who supports these teams is being arranged.
3	Workflow Integration	3 ■■■■ <i>AI-proposed</i>	Builder claim	Recordings move in and out as Scripture Burrito files, in both directions, with verse alignment preserved, and USFM can be imported. In practice the connection we can see is to the builder's own companion app; whether it works as smoothly with anyone else's software is untested. The top of the scale would need the hand-over to happen automatically rather than by exporting and importing.
4	User Reach & Adoption <i>Needs real users</i>	3 ■■■■ <i>AI-proposed</i>	Builder claim	More than forty teams across several clusters have used Scribe from start to finish, all within the builder's own translation programme. Use elsewhere has been mentioned to us — a project in Africa, and a sign-language team in the Solomon Islands — but we have not been able to confirm either. How many are working with it now is not known. Would change it: current active-team count and any verified use outside BCS
5	Ease of Use & Learning Curve	— ■■■■	Pending	How much training teams need, and where translators get stuck, can only come from the teams and the facilitator who supports them.
6	Project Implementation <i>Needs real users</i>	4 ■■■■ <i>AI-proposed</i>	Builder claim	Forty completed oral Bibles were drafted in Scribe, which is well past the five projects the top of the scale asks for. That scale asks about projects relying on it now, and that number is not known. The forty came from several parallel clusters in one programme rather than forty separate decisions to adopt it.
7	Implementation Horizon	2 ■■■■ <i>AI-proposed, contested</i>	Documentation	Version 1.2.0 was released on 2 July 2026, adding video pause and resume, video import and audio joining. The builder describes the tool as finished and maintained, with no new features planned. It sits awkwardly on this scale. The wording for 3 asks for core features ready for use and a clear long-term roadmap: Scribe has the first but not the second. The wording that describes having no roadmap is the one for 1, which also describes unclear usefulness today, and that does not fit a tool in active use. The wording for 2 describes limited functionality with a roadmap that lacks a timeline, which fits neither. No description fits a tool that is complete and maintained but no longer growing, which is why this row is marked contested. The full wording is in the appendix. Would change it: a rubric anchor for finished-and-maintained tools
8	Workflow Stage Fit	4 ■■■■ <i>AI-proposed, contested</i>	Builder claim	Scribe covers drafting and peer checking with re-recording, two stages that flow into each other, and hands work onward without losing verse alignment. Reviewers may prefer 3, since these are two of the six stages in the builder's own pipeline and the wording does not say whether covering multiple stages means a count or a proportion.

9	Post-Implementation Support <i>Needs real users</i>	3 ■■■■ <i>AI-proposed</i>	Builder claim	There is a direct support address, staff at the builder who answer it, a documentation site and a community forum. The top of the scale also asks for ongoing training and regular updates, and the tool is no longer being updated. Whom teams actually contact, and how quickly problems get solved, will come from the interviews.
10	Retention & Sustained Impact <i>Needs real users</i>	— ■■■■	Pending	Which teams are still working in it, which stopped and why, can only come from the facilitator and the teams. The forty published Bibles show that projects finished, which is not the same as teams staying with the tool.
11	Cost & Financial Sustainability <i>Decisive</i>	3 ■■■■ <i>AI-proposed</i>	Builder claim	If Scribe disappeared in two years, a team's recordings would remain: the tool is offline by default, the files are on the team's own machine, and the format it writes is an open one that other software can read. It is free and open source, so the code outlives whatever happens to its builder, and one team could in principle keep maintaining it. The tool is finished rather than growing, which cuts both ways — less to maintain, but also no active development to lose. The dependency is not Scribe itself but the online companion the builder's own teams use for checking afterwards, which is a service rather than a file on a laptop. Would change it: a stated funding horizon, and clarity on what the online companion costs a team
12	Data Governance & Sovereignty <i>Decisive</i>	3 ■■■■ <i>AI-proposed</i>	Builder claim	Everything stays on the team's own machine unless they choose to sync, either to a shared repository — which requires creating an account and accepting an open licence first — or to a server they run themselves. For the wider workflow: the companion app used later for checking holds team audio and text on the builder's servers, and we do not know its retention terms.
13	Explainability & Human Steering	N/A ■■■■	Not applicable	There is no AI in Scribe, so this question does not apply to it. The workflow Scribe feeds does depend on AI: the companion app transcribes the recordings and generates speech from the checked text.
14	Cultural & Safety Alignment	3 ■■■■ <i>AI-proposed</i>	Builder claim	Scribe brings no interpretive content of its own. It fetches whichever published open-licensed resources a team chooses and lets teams load their own material, so what a translator sees beside their recording is the team's choice. Which resources, in which language, were actually open while these clusters were drafting is a question for the interviews.

Still open

- No user evidence yet. A facilitator call is being scheduled; team interviews depend on it, and language and distance are barriers.
- How many teams are drafting now, versus the 40+ that have finished; how many clusters are still in progress.
- Whether teams sync to Door43, and whether they understood the licence they accepted.
- BCS's funding horizon for Scribe, and the terms of use and retention policy of the companion app.
- Interface languages: the builder names six, the documentation five. Both recorded.

Builder's note. The builder has not yet reviewed this page.

Where to find it

scribe.bible · source code at
github.com/bible-technology/scribe-scripture-editor ·
support by email to scribe@bridgeconn.com and
through the Door43 community forum · available for
Windows, Linux and macOS, and in a browser.

The user it assumes

An oral translation team with a laptop; a facilitator or support person for setup, to be confirmed; an interface language among English, Hindi, Telugu, Kannada, French and Russian; and a route out for checking and publication, which in BCS's programme is AI OBT, an online tool needing an account.

How to read this

A score is chosen by reading the four descriptions in the appendix and picking the one that fits the evidence.

AI-proposed means a draft score, prepared from the evidence on file and still to be reviewed by two people.

Contested means the evidence points both ways, or the rubric's own wording does not settle the case — the reasoning beside it is worth reading. Figures the builder gave us are labelled as theirs. Scores are never added together or averaged.

How this page was put together. Findings were gathered by the evaluation team through a structured interview with the builder and a reading of the documentation and code repository. The scores are drafts prepared with AI assistance from that evidence and are still to be reviewed by two people.

Appendix: the full rubric

The wording every score is measured against. A score is chosen by reading these four descriptions and picking the one that fits what the evidence shows — not by impression, and not on a curve. ● decisive criterion · ● needs real users (Tier 1 only).

Criterion	1 — Unsatisfactory	2 — Developing	3 — Proficient	4 — Exemplary
1 Problem Alignment <i>What problem does it solve?</i> ● decisive	Solves an ill-defined or non-existent problem; no alignment with target audience needs.	Addresses a minor issue, but core pain points remain unresolved.	Solves a clear, well-defined problem relevant to field translation teams.	Solves a high-priority, critical problem with significant, measurable impact.
2 Setup Complexity <i>How easy is it to set up?</i>	Extremely complex; requires specialized technical skills, custom code, or extensive workarounds.	Difficult setup; requires moderate technical support or multi-step configuration.	Straightforward setup; standard user can complete configuration with documentation.	Turnkey, plug-and-play setup; minimal to no technical intervention required.
3 Workflow Integration <i>Can this tool be easily integrated into the workflow?</i>	Incompatible; disrupts existing workflows and creates significant friction or duplicate work.	Integrates with difficulty; requires manual data re-entry or awkward workarounds.	Integrates smoothly into standard field tools and processes with minimal friction.	Seamless native integration; fully automated data flow between tools and platforms.
4 User Reach & Adoption <i>How many people are using it, and who are they?</i> ● needs real users	No reachable active users; purely theoretical or tested only in demo settings.	Very few active users (1–5); restricted to a single closed group.	Moderate active user base across target user roles.	Large, established active user base across multiple field teams and roles.
5 Ease of Use & Learning Curve <i>How intuitive is the tool to use daily?</i>	Unintuitive interface; steep learning curve requiring extensive ongoing training or technical intervention.	Somewhat complex interface; requires moderate user training and frequent reliance on support manuals.	User-friendly and intuitive for everyday translation work; minimal training required for field teams.	Highly intuitive and accessible; seamless navigation with virtually no learning curve for non-technical users.
6 Project Implementation <i>How many projects have implemented this tool?</i> ● needs real users	0 real-world projects implemented.	1 active pilot project.	2–4 active field projects currently using the tool.	5+ real-world translation projects actively relying on the tool.
7 Implementation Horizon <i>What can be implemented now, and what is promised later?</i>	Unclear immediate utility and no long-term roadmap.	Limited current functionality; long-term roadmap exists but lacks timeline or feasibility.	Core features ready for immediate deploy; clear long-term feature roadmap in place.	Fully functional immediate capabilities paired with a realistic, high-value long-term plan.

Criterion	1 — Unsatisfactory	2 — Developing	3 — Proficient	4 — Exemplary
8 Workflow Stage Fit <i>Where in its domain's workflow does the tool fit, and how much of that workflow does it cover?</i>	Cannot be placed at any recognised stage of its declared domain's workflow, or the fit is so narrow or ill-defined that a team could not say when in their process to use it.	Usable at a single stage, but with significant constraints or gaps; the stage's core task is only partly addressed.	Fully functional at one or more core stages of its domain's workflow, with a clear hand-off into and out of the tool.	Covers multiple stages of its domain's workflow with continuity between them — outputs of one stage feed the next without re-entry.
9 Post-Implementation Support <i>What support is available after a team adopts it?</i> ● needs real users	No post-launch support, contact points, or documentation available.	Minimal support; limited to unmaintained community forums or slow response times.	Reliable support available via direct contact channels with clear user documentation.	Comprehensive support including direct helpdesk, ongoing training, and regular updates.
10 Retention & Sustained Impact <i>What difference is it making on the ground, over time?</i> ● needs real users	Tool is abandoned shortly after launch; zero sustained adoption or measurable impact on ground operations.	High drop-off rate after initial trial; minimal lasting change in team productivity or workflows over 6 months.	Consistent usage sustained after 6+ months; clear, positive shift in field team efficiency or output quality.	Entrenched as an essential tool with near-total user retention; measurable, long-term transformational impact on ground work.
11 Cost & Financial Sustainability <i>Is the pricing model viable and sustainable long-term?</i> ● decisive	Cost-prohibitive or opaque pricing; reliant on unsustainable short-term subsidies with high financial risk.	High operational or licensing costs; partially sustainable but creates budget strain over time.	Predictable, fair pricing structure supported by a clear, long-term financial plan.	Free or open-access, or highly cost-effective; backed by a fully sustainable funding and maintenance model.
12 Data Governance & Sovereignty <i>How are data ownership, privacy, and security handled?</i> ● decisive	No data privacy controls; vendor retains unrestricted ownership or rights over sensitive translation content.	Basic security measures, but vague data retention, export, and regional privacy compliance policies.	Clear data ownership rights retained by the translation team, with standard privacy and export protocols.	Complete data sovereignty guaranteed; robust encryption, strict privacy safeguards, and full local data control.
13 Explainability & Human Steering <i>Can users understand and control why the AI gives an answer?</i>	Opaque “black box” system; no rationale provided for outputs, and users cannot adjust, override, or guide model decision-making.	Minimal rationale or source tracking; limited ability to tweak prompt settings or influence underlying output logic.	Clear output explanations or source references provided; users can easily adjust steering parameters and override AI outputs.	Fully transparent logic, traceable citations, and granular human controls to tune, override, and direct AI decision-making at every stage.
14 Cultural & Safety Alignment <i>Are outputs culturally accurate, ethically aligned, and free of regional bias?</i>	Outputs display significant regional bias, offensive or harmful content, or complete disregard for local cultural context.	Inconsistent safety controls; occasional cultural inaccuracies, subtle regional biases, or weak ethical boundary safeguards.	Generally reliable safety and ethical boundaries; culturally appropriate outputs with minimal regional bias.	Rigorously calibrated for localized ethics, nuance, and safety; zero regional bias with complete context awareness and customization.