

How to read this

Each tool has two pages. The **Summary page** gives the short version: what the tool is, who has told us what, and where the evidence stands. The **Findings & Evidence** page sets out all fourteen criteria, the score each has been given so far, and what that score rests on. Read the basis column rather than the number: the number is shorthand, the basis is the finding.

Where a score comes from

Every criterion is scored 1 to 4 against wording agreed in advance, so that a 3 means the same thing for every tool. Scores are never added up or averaged, and there is no overall mark. A tool that scores 2 on something your organisation does not care about may still be the right tool for you.

agreed	Reviewers scored it independently, then compared scores and agreed the result on the record. Where they had differed, they settled it by discussion, not by averaging or a vote.
provisional	Every reviewer who scored it gave the same number, but they have not yet discussed it. None appear in this edition yet.
reviewers differ	Reviewers scored it independently and landed on different numbers. Every score given is shown; the bar on the chart shows the lowest. None appear in this edition yet.
AI-proposed	A draft score with its reasoning, prepared with AI assistance from the evidence we gathered. It is not a human judgement. None appear in this edition yet.
contested	The evidence splits, or our own wording does not settle the case. Worth reading the basis.

Where the evidence comes from

Every score carries a tag saying what it rests on. The tags are not interchangeable: what a translator reports from a live project counts for more than what a product page says, and both are printed so you can weigh them yourself.

User-reported	A translator or team using the tool in a live project told us directly, through a survey or an interview we ran.
Facilitator-reported	A facilitator, coordinator or support worker told us about the teams they serve. Close to the work, but second-hand.
Our testing	We used the tool ourselves against a fixed set of questions and recorded what happened.
Documentation	A published page, release note, policy or repository, with the date we read it.
Builder claim	The people who make the tool told us. Printed as theirs, not as established fact.
Pending	Asked for but not yet in hand. The page says what was requested and when.
N/A	The criterion does not apply — for example, explainability for a tool with no AI in it.

The three criteria we treat as decisive

Eleven of the fourteen criteria describe how well a tool does something. Three describe whether it is safe to rely on at all, and a score of 1 on any of them is worth more of your attention than a handful of 2s elsewhere. They are marked in red on the Findings & Evidence page. Nothing here removes a tool from anyone’s list — that decision belongs to your organisation — but these are the three we would want answered before committing a project to a tool.

- 1. Problem Alignment** — whether the tool solves a problem translation teams actually have. A tool that is impressive but answers a question nobody asked will not survive contact with a project, however well it scores elsewhere.
- 11. Cost & Financial Sustainability** — whether a team can still afford it in three years. Adoption is not a purchase but a dependency: a tool whose funding quietly ends leaves a project stranded mid-book.
- 12. Data Governance & Sovereignty** — who holds the translation text, on what terms, and who else can see it. Unpublished Scripture is a community’s own work and sometimes a safety matter, and unlike the other criteria a mistake here cannot be undone later.

The two questions we ask before a tool goes on this list

Can we assess it? That is: is it in scope, and can we find out enough about it? A tool qualifies if it serves Bible translation or related ministry and if we can gather enough information to answer the fourteen questions — from the people using it where they exist, and otherwise from documentation, from the people who built it and from our own use. A tool we cannot find out enough about is named in the wider catalogue and marked as not assessed, rather than being scored on guesswork.

Is it worth your time here? That is: among the tools we could assess, we gave priority to those that address a problem organisations in this room actually have, that more than one team has used, or that stand for a choice many of you are weighing this year. This is a decision about what fits a short session, not a ranking, and a tool being absent says nothing against it.

The fourteen criteria

Every tool is asked the same fourteen questions. ● marks the three we treat as decisive; ● marks the four that need real users before they can be answered. The full wording for each score, 1 to 4, is in the appendix at the back.

1. **Problem Alignment** ● — What problem does it solve?
2. **Setup Complexity** — How easy is it to set up?
3. **Workflow Integration** — Can this tool be easily integrated into the workflow?
4. **User Reach & Adoption** ● — How many people are using it, and who are they?
5. **Ease of Use & Learning Curve** — How intuitive is the tool to use daily?
6. **Project Implementation** ● — How many projects have implemented this tool?
7. **Implementation Horizon** — What can be implemented now, and what is promised later?
8. **Workflow Stage Fit** — Where in its domain's workflow does

the tool fit, and how much of that workflow does it cover?

9. **Post-Implementation Support** ● — What support is available after a team adopts it?
10. **Retention & Sustained Impact** ● — What difference is it making on the ground, over time?
11. **Cost & Financial Sustainability** ● — Is the pricing model viable and sustainable long-term?
12. **Data Governance & Sovereignty** ● — How are data ownership, privacy, and security handled?
13. **Explainability & Human Steering** — Can users understand and control why the AI gives an answer?
14. **Cultural & Safety Alignment** — Are outputs culturally accurate, ethically aligned, and free of regional bias?

Why some pages leave four criteria unscored

Four of the fourteen can only be answered by people using the tool in real work: **4 User Reach & Adoption**, **6 Project Implementation**, **9 Post-Implementation Support** and **10 Retention & Sustained Impact**. Where we could reach such users, a tool is assessed against all fourteen — what we call **Tier 1**. Where we could not, it is assessed against the other ten only, which we call **Tier 2**, and those four are left deliberately unscored rather than guessed at.

A Tier 2 page is not a worse tool or a lesser assessment. It usually means the tool is new, or that its users were beyond our reach in the time we had. Newer tools cannot show a user base until someone adopts them, so treating the absence of one as a mark against a tool would make it impossible for any new tool to be adopted at all. Where a tool is at Tier 2, the page says what we would need in order to answer the remaining four.

Please read this before using the scores. Every score in this document was agreed by our reviewers, who scored independently and then settled their scores together on the record, by discussion rather than by averaging or a vote. The evidence behind every score — the surveys, interviews, correspondence with builders, trial records and our own testing — is the work of our team, and each page says where every finding came from. Please treat the numbers as a starting point for your own judgement, not as settled findings.



Scripture Forge

SIL · scriptureforge.org · Free, open source (MIT) · Web, linked to Paratext · Uses AI to draft

Executive summary

Scripture Forge generates a first-pass draft of a Bible book in a team’s own language, using the team’s completed work as the model, and delivers it into their Paratext project. It is used in several ways: drafting New Testament books, drafting incrementally as a translation grows, back translation, adaptation projects, and community checking. Its most common use is drafting the Old Testament once a team’s New Testament is finished, and that is the use this assessment mostly reflects: twelve of the fifteen teams we heard from had a completed New Testament and were drafting the Old.

This assessment draws on written responses from fifteen people across seven organisations, gathered between 18 and 28 September 2026 and covering nineteen named language projects in India and Nepal, with three respondents who did not name theirs; a reading of Scripture Forge’s published terms of use and privacy policy; and figures supplied by SIL. Thirteen of the fifteen are still using it.

Most teams say it solved the problem they came with, and ten have been using it for over a year. Two things temper that. Drafts are rewritten heavily — most teams keep between one and eight verses in ten with only small changes, and four teams say drafting is now slower than it was before. And nine of the fifteen have seen wording carried over from the source that does not suit their language, which for one team is a structural mismatch between an Indo-Aryan source and a Dravidian target that they correct by hand, book after book.

The main open question is what happens to the text a team sends for drafting. The terms of use predate the AI feature entirely and say nothing about it, so only SIL can say where that text is processed and whether it is used to improve their models.

What it is for, and where it fits

Written translation: drafting, community checking and back translation

Scored against the written translation workflow

Our connections to this tool

We support several of these teams in their use of Scripture Forge. We are not connected to the people who build the tool, but we are close enough to these teams to expect that some will have softened what they told us.

How much evidence this rests on

Fifteen users, seven organisations, four evidence sources



Responses gathered 18–28 September 2026, from teams in India and Nepal

Who built it

Built by SIL. The figures on this page are SIL’s own.

Not known: how many projects and people are active now, and what happens to text sent for drafting.

Figures the builder gave us

Supplied by the people who make the tool. We print them as given and have not verified them independently.

BUILDER-REPORTED

274

projects in India the builder counts as using it

Includes some back-translation projects. It does not say how many are active now. The builder also gave a lower figure, 121, for projects actively using it; the difference between the two was not explained.

BUILDER-REPORTED

234

people with a role on those projects

Each person counted once. Many hold roles on several projects — administrators and support staff especially — so this counts people, not project places. It does not say how many used it recently.

BUILDER-REPORTED


852

book drafts generated since September 2025

Since September 2025, the earliest point from which the builder can separate real projects from test ones. It counts drafts produced, including back translations, not drafts a team went on to use.

All fourteen questions at a glance

One bar per question, numbered to match the rows overleaf. The taller the bar, the higher the score out of four; the colour shows what kind of evidence that score rests on. Outlined bars are questions we have not been able to answer yet.



User-reported

Documentation

Pending

Red numbers: veto criteria

Where the evidence came from

User-reported	11
Documentation	3

Before scoring

Can we assess it? Yes. In scope; assessable from documentation and builder data; named, reachable users heard from via the form, so Tier 1.

Is it worth your time here? Yes. Drafting pace is a problem every organisation has; evidence exists from more than one team; AI drafting is the category decision most of the audience faces this year.

What the fifteen users told us

- Twelve said Scripture Forge solved the problem they adopted it for, one said partly, and two said it did not.
- Ten find drafting faster, three of them much faster. Four find it slower, and two of those four still use it regularly — the cost is worth paying for them, but it is a cost.
- Someone other than the team did the setup in eleven of the fifteen cases. Eight teams needed several training sessions or ongoing technical help.
- Drafts are rewritten heavily. Out of every ten verses, two teams keep one or two, five keep three to five, four keep six to eight, and one keeps none.
- Two teams have stopped, both within three months, and both because the draft was not worth the editing.
- Nine of the fifteen have seen wording carried over from the source that does not suit their language — a dragon rendered as pigs, borrowed words for concepts with no precedent, and for one team a systematic word-order mismatch between their source language and their own.
- Eight know where their translation text is stored and who can reach it; four know roughly; two do not know. What they accepted are SIL's general Language Software policies,

which state that they were last updated in September 2019 and say nothing about AI drafting.

Findings & Evidence

Each score sits against a written anchor (1 unsatisfactory, 2 developing, 3 proficient, 4 exemplary). The basis column says what the score rests on and what would change it. Each score below has been agreed by two reviewers working independently.

	Criterion	Score	Evidence	Basis
1	Problem Alignment <i>Decisive</i>	3 ■■■■ agreed	User-reported	Teams came to it for different reasons: several wanted drafting to go faster, one had no experienced translator available, two wanted more natural language and better key terms, one wanted Old Testament drafts generated from a New Testament already finished, and one wanted to speed up back translation, which had been holding up their targets. Twelve of the fifteen said it solved the problem they came with, one said partly, and two said it did not. On speed the picture is mixed and has become more so: three find drafting much faster, seven somewhat faster, one about the same, and four slower. The problem is real and the tool addresses it for most teams, but a quarter of them are working more slowly than before.
2	Setup Complexity	3 ■■■■ agreed	User-reported	Someone other than the team did the setup in eleven of the fifteen cases — their own organisation's support staff in seven, a partner's in one, the Scripture Forge team in two, and a named individual in one. Four teams had a member with administrator rights do it themselves. From deciding to use it to a first usable draft took under a week for four teams, one to four weeks for three, and one to three months for four; four could not say, mostly because someone else did the work. One team lost time to an incorrect language code in Paratext, and another to markers in their Paratext project.
3	Workflow Integration	3 ■■■■ agreed	User-reported	Eleven of the fifteen chose direct sync into Paratext. Most of the rest copy and paste, and one mostly uses direct sync but into a separate AI project, mainly for books already translated by hand, so that the AI draft can be compared and improved. The friction teams named is specific and recurring: one re-uploads a spreadsheet of proper names with every draft request and must delete the previous one first; one repeats work on some passages; one runs the generation command many times; and one describes converting the AI draft for each book as a step of its own.
4	User Reach & Adoption <i>Needs real users</i>	3 ■■■■ agreed	User-reported	Fifteen people answered, from seven organisations, covering nineteen named language projects — one respondent alone speaks for seven — and three who did not name theirs. Most are in India; at least one is in Nepal. That is a substantial spread of teams and roles for a tool of this kind. Against it, the builder counts 274 projects in India using Scripture Forge, some of them back translations, with 234 people holding roles on them; many hold roles on several projects. The builder also gave a lower figure, 121, for projects actively using it, and the difference between the two was not explained. Neither figure says how many teams are active now.
5	Ease of Use & Learning Curve	3 ■■■■ agreed	User-reported	Asked how much training they needed for day-to-day use, one team said none, six a short walkthrough, four several sessions, and four ongoing help from someone technical — so a little over half needed sustained support rather than a one-off introduction. The obstacles named were the team's willingness to use it, an initial language-code problem, marker problems in Paratext, slow internet in the field, and the labour of correcting word order by hand where the source language and the target language differ in structure. Four teams reported no obstacle.
6	Project Implementation <i>Needs real users</i>	4 ■■■■ agreed	User-reported	Nineteen named language projects across seven organisations use it, and three respondents did not name theirs: eleven of the fifteen respondents use it regularly, two occasionally, and two have stopped. Ten have been using it for more than a year, and seven have drafted more than ten books each. That is comfortably past the top of the scale, which asks for five or more projects actively relying on the tool.

7	Implementation Horizon	4 ■■■■ agreed	Documentation	The core features work today and are in daily use across many projects, and the tool is actively maintained. We have not found a published roadmap with dates, which is what the top of the scale asks for.
8	Workflow Stage Fit	3 ■■■■ agreed	User-reported	All fifteen teams use it for drafting, seven also generate back-translation drafts, and two use it for community checking or review. None edits inside Scripture Forge itself. Because drafts sync directly into Paratext for most teams, the drafting stage hands over cleanly to the team's own working project. The checking stages exist but are barely used, so continuity beyond drafting is still not demonstrated.
9	Post-Implementation Support Needs real users	3 ■■■■ agreed	User-reported	Support reaches people through several routes and generally works. Teams described contacting their own organisation's IT staff and getting a fix within a day; the Scripture Forge team by email, without much delay; and named individuals directly, including over WhatsApp. One issue took two days to resolve because it had to be handled remotely. Two respondents had nothing to report.
10	Retention & Sustained Impact Needs real users	3 ■■■■ agreed	User-reported	Thirteen of the fifteen are still using it, eleven of them regularly, and ten have been using it for over a year. Asked how much of an AI draft survives with only small changes, two teams keep one or two verses in ten, five keep three to five, four keep six to eight, one keeps none, and one could not say. Two teams have stopped, both within three months, and both gave the same reason: the draft quality was not worth the editing. One of those respondents added that the team's own willingness played a part; the other said quality was low for one of their two languages. Usage is sustained and drafts are heavily rewritten, and because we reached teams through the people who support them, teams who quietly gave up are still likely to be under-represented.
11	Cost & Financial Sustainability Decisive	4 ■■■■ agreed	Documentation	If Scripture Forge disappeared in two years, a team would lose its drafting head start and go back to drafting by hand — but nothing already produced would be lost, because every draft lands in the team's own Paratext project and stays there. Teams are not locked in. The tool is free and open source, so the code would survive its funder; what would not survive is the service that generates the drafts, which costs its builder money every time a team uses it, and which no team pays for. No end date has been published, and none has been suggested to us.
12	Data Governance & Sovereignty Decisive	2 ■■■■ agreed	Documentation	We read the terms of use and the privacy policy on 23 September 2026. Both are SIL's general Language Software policies rather than anything specific to Scripture Forge, and both state that they were last updated on 30 September 2019; neither mentions AI, models, training or draft generation, so a team accepting them is not agreeing to anything specific about AI drafting. Ownership is clear: sharing data does not change who owns it. But uploading grants SIL a wide licence to host, copy, modify, make derivative works from, publish and distribute the text, for as long as it is hosted there. Translation text counts as Content rather than personal information, so the protections the privacy policy describes for personal data do not apply to it, and data from different users and services may be combined. Nothing is said about how long text is kept, where it is processed, or how to delete it from Scripture Forge, and no regional privacy law is mentioned. Disputes are handled in Texas, and an organisation using the service accepts the terms on behalf of its teams. Among the users we heard from, eight know where their text is stored and who can reach it, four know roughly, and two do not know.
13	Explainability & Human Steering	2 ■■■■ agreed	User-reported	Asked whether they can tell why a draft chose a particular wording, three said usually, seven sometimes, one never, and four did not answer. Drafts arrive without any explanation attached. The way

			teams steer the output is indirect: feeding corrected text back in so later books improve, supplying a spreadsheet of proper names, or changing which translation is used as the source.
14	Cultural & Safety Alignment	2■■■■ agreed User-reported	Nine of the fifteen had seen wording or ideas carried over from the source that did not fit their language or culture; three had not noticed any, two said no, and one did not answer. The examples are specific: a dragon rendered as pigs; borrowed words appearing where a concept had no precedent in the data; one team's back translation of Isaiah 34:15, which the respondent called "funny" — the wild goats calling their friends' goats; and vocabulary drawn from books the team had not yet translated. One team described the underlying problem structurally — their source text is an Indo-Aryan language and their target is Dravidian, so word order and clause order come out wrong and have to be corrected by hand, and they cannot substitute a Dravidian source text because it would diverge from the regional Bible their community expects. These look like limits of the material the model has to work from rather than a systematic bias.

Still open

- What happens to text sent for drafting: whether it trains or improves models, where it is processed, how long it is kept. The terms of use and privacy policy, both dated 30 September 2019, predate the feature and are silent on it. Only the builder can answer.
- How many of the 234 people with a role on these projects have used it recently, and how many projects are active now.
- Why four teams find drafting slower while still choosing to use it, and what separates them from the teams who find it much faster.
- Whether the teams who stopped are typical of others we did not reach. Both stopped teams came to us through the people who support them, as everyone in this survey did.
- How much of a draft actually survives, measured rather than estimated — by comparing drafts with the finished text for a few teams.

Where to find it

scriptureforge.org · help and release notes at help.scriptureforge.org · terms at scriptureforge.org/terms and privacy at scriptureforge.org/privacy · generating a draft needs a Paratext project with roughly a complete New Testament across a range of genres.

The user it assumes

NT complete and OT in progress; Paratext linked; a technically capable member or support staff; reliable internet; source texts licensed for training; someone with Paratext administrator rights; a workflow that accepts post-editing a draft zero; and an organisation willing to accept SIL's generic language-software terms — under Texas jurisdiction, with the organisation indemnifying SIL — on behalf of its teams.

How to read this

A score is chosen by reading the four descriptions in the appendix and picking the one that fits the evidence. **Contested** means the evidence points both ways, or the rubric's own wording does not settle the case — the reasoning beside it is worth reading. Figures the builder gave us are labelled as theirs. Scores are never added together or averaged.

How this page was put together. Findings were gathered by the evaluation team through a user survey, correspondence with SIL and a reading of the published terms. The scores are drafts prepared with AI assistance from that evidence and are still to be reviewed by two people; they are marked accordingly in the table above.

Appendix: the full rubric

The wording every score is measured against. A score is chosen by reading these four descriptions and picking the one that fits what the evidence shows — not by impression, and not on a curve. ● decisive criterion · ● needs real users (Tier 1 only).

Criterion	1 — Unsatisfactory	2 — Developing	3 — Proficient	4 — Exemplary
1 Problem Alignment <i>What problem does it solve?</i> ● decisive	Solves an ill-defined or non-existent problem; no alignment with target audience needs.	Addresses a minor issue, but core pain points remain unresolved.	Solves a clear, well-defined problem relevant to field translation teams.	Solves a high-priority, critical problem with significant, measurable impact.
2 Setup Complexity <i>How easy is it to set up?</i>	Extremely complex; requires specialized technical skills, custom code, or extensive workarounds.	Difficult setup; requires moderate technical support or multi-step configuration.	Straightforward setup; standard user can complete configuration with documentation.	Turnkey, plug-and-play setup; minimal to no technical intervention required.
3 Workflow Integration <i>Can this tool be easily integrated into the workflow?</i>	Incompatible; disrupts existing workflows and creates significant friction or duplicate work.	Integrates with difficulty; requires manual data re-entry or awkward workarounds.	Integrates smoothly into standard field tools and processes with minimal friction.	Seamless native integration; fully automated data flow between tools and platforms.
4 User Reach & Adoption <i>How many people are using it, and who are they?</i> ● needs real users	No reachable active users; purely theoretical or tested only in demo settings.	Very few active users (1–5); restricted to a single closed group.	Moderate active user base across target user roles.	Large, established active user base across multiple field teams and roles.
5 Ease of Use & Learning Curve <i>How intuitive is the tool to use daily?</i>	Unintuitive interface; steep learning curve requiring extensive ongoing training or technical intervention.	Somewhat complex interface; requires moderate user training and frequent reliance on support manuals.	User-friendly and intuitive for everyday translation work; minimal training required for field teams.	Highly intuitive and accessible; seamless navigation with virtually no learning curve for non-technical users.
6 Project Implementation <i>How many projects have implemented this tool?</i> ● needs real users	0 real-world projects implemented.	1 active pilot project.	2–4 active field projects currently using the tool.	5+ real-world translation projects actively relying on the tool.
7 Implementation Horizon <i>What can be implemented now, and what is promised later?</i>	Unclear immediate utility and no long-term roadmap.	Limited current functionality; long-term roadmap exists but lacks timeline or feasibility.	Core features ready for immediate deploy; clear long-term feature roadmap in place.	Fully functional immediate capabilities paired with a realistic, high-value long-term plan.

Criterion	1 — Unsatisfactory	2 — Developing	3 — Proficient	4 — Exemplary
8 Workflow Stage Fit <i>Where in its domain's workflow does the tool fit, and how much of that workflow does it cover?</i>	Cannot be placed at any recognised stage of its declared domain's workflow, or the fit is so narrow or ill-defined that a team could not say when in their process to use it.	Usable at a single stage, but with significant constraints or gaps; the stage's core task is only partly addressed.	Fully functional at one or more core stages of its domain's workflow, with a clear hand-off into and out of the tool.	Covers multiple stages of its domain's workflow with continuity between them — outputs of one stage feed the next without re-entry.
9 Post-Implementation Support <i>What support is available after a team adopts it?</i> ● needs real users	No post-launch support, contact points, or documentation available.	Minimal support; limited to unmaintained community forums or slow response times.	Reliable support available via direct contact channels with clear user documentation.	Comprehensive support including direct helpdesk, ongoing training, and regular updates.
10 Retention & Sustained Impact <i>What difference is it making on the ground, over time?</i> ● needs real users	Tool is abandoned shortly after launch; zero sustained adoption or measurable impact on ground operations.	High drop-off rate after initial trial; minimal lasting change in team productivity or workflows over 6 months.	Consistent usage sustained after 6+ months; clear, positive shift in field team efficiency or output quality.	Entrenched as an essential tool with near-total user retention; measurable, long-term transformational impact on ground work.
11 Cost & Financial Sustainability <i>Is the pricing model viable and sustainable long-term?</i> ● decisive	Cost-prohibitive or opaque pricing; reliant on unsustainable short-term subsidies with high financial risk.	High operational or licensing costs; partially sustainable but creates budget strain over time.	Predictable, fair pricing structure supported by a clear, long-term financial plan.	Free or open-access, or highly cost-effective; backed by a fully sustainable funding and maintenance model.
12 Data Governance & Sovereignty <i>How are data ownership, privacy, and security handled?</i> ● decisive	No data privacy controls; vendor retains unrestricted ownership or rights over sensitive translation content.	Basic security measures, but vague data retention, export, and regional privacy compliance policies.	Clear data ownership rights retained by the translation team, with standard privacy and export protocols.	Complete data sovereignty guaranteed; robust encryption, strict privacy safeguards, and full local data control.
13 Explainability & Human Steering <i>Can users understand and control why the AI gives an answer?</i>	Opaque “black box” system; no rationale provided for outputs, and users cannot adjust, override, or guide model decision-making.	Minimal rationale or source tracking; limited ability to tweak prompt settings or influence underlying output logic.	Clear output explanations or source references provided; users can easily adjust steering parameters and override AI outputs.	Fully transparent logic, traceable citations, and granular human controls to tune, override, and direct AI decision-making at every stage.
14 Cultural & Safety Alignment <i>Are outputs culturally accurate, ethically aligned, and free of regional bias?</i>	Outputs display significant regional bias, offensive or harmful content, or complete disregard for local cultural context.	Inconsistent safety controls; occasional cultural inaccuracies, subtle regional biases, or weak ethical boundary safeguards.	Generally reliable safety and ethical boundaries; culturally appropriate outputs with minimal regional bias.	Rigorously calibrated for localized ethics, nuance, and safety; zero regional bias with complete context awareness and customization.